

ARJUN MEHTA

Final-year B.Tech CSE | arjun.mehta@example.com | github.com/arjunm

SKILLS

Python, Go, PyTorch, FastAPI, PostgreSQL, Redis, Kafka, Docker, GCP

PROJECTS

DataMesh - Streaming feature pipeline

Built an async ingestion proxy in Python handling roughly 2000 requests per second across 12 upstream sources. Used one Kafka topic per source and deduplicated downstream with a bloom filter before writing to the feature store. Cut end-to-end feature staleness from 40s to under 5s.

RecSys-Lite - Retrieval and ranking service

Two-stage recommender: ANN candidate retrieval over FAISS, then a gradient-boosted ranker. Served with FastAPI behind Redis caching. Improved click-through rate by 18% in an offline replay evaluation.

ShardKV - Distributed key-value store (course project)

Implemented Raft consensus in Go with log compaction and snapshotting. Supported dynamic shard rebalancing across replica groups.

EXPERIENCE

Software Engineering Intern, Netlyx Systems (Summer 2025)

Migrated a monolithic billing job to a queue-driven worker pool, reducing nightly batch runtime from 6 hours to 90 minutes. Added Prometheus metrics and on-call runbooks for the new pipeline.

EDUCATION

B.Tech Computer Science and Engineering, 2026. CGPA 8.7/10.

Relevant coursework: Distributed Systems, Machine Learning, Databases.